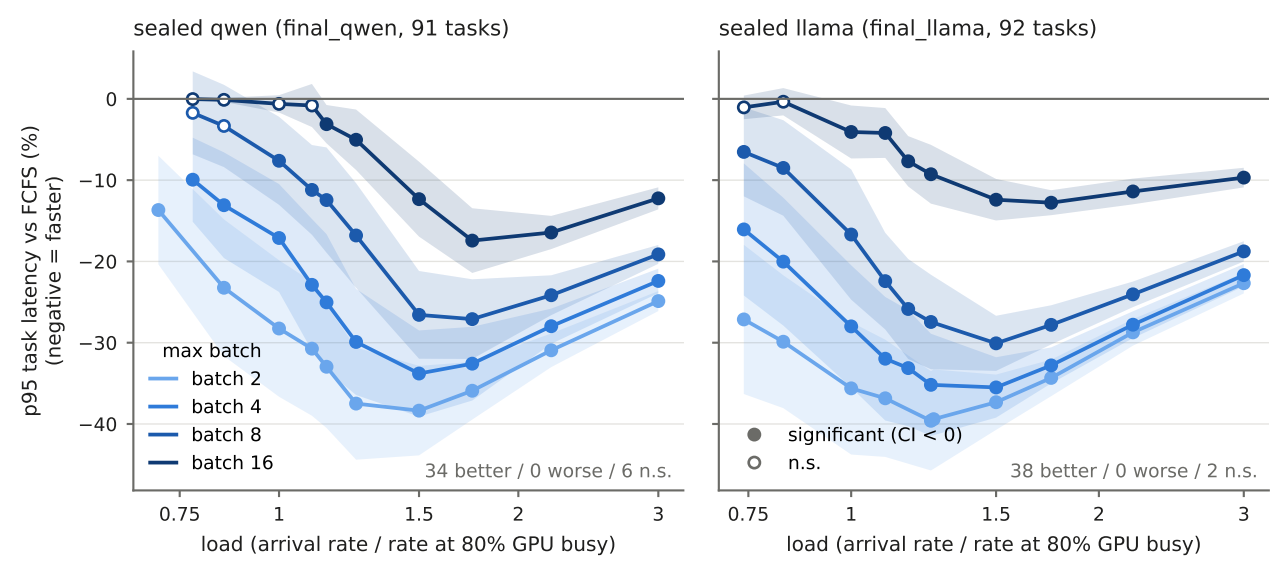
sealed qwen (final_qwen, 91 tasks)
sealed llama (final_llama, 92 tasks)
p95 task latency vs FCFS (%)
(negative = faster)
0
−10
−20
−30
−40
max batch
batch 2
batch 4
batch 8
batch 16
34 better / 0 worse / 6 n.s.
38 better / 0 worse / 2 n.s.
significant (CI < 0)
n.s.
0.75
1
1.5
2
3
load (arrival rate / rate at 80% GPU busy)